

DATA621-FinalProject-SmoothOperators

Rob Hodde, Matt Farris, Jeffrey Burmood, Bin Lin

5/26/2017

Problem Description

Our final project will explore, analyze and model a data set containing information on approximately 5,000 movies. The dataset contains movie data extracted from the IMDB website and is available on Kaggle.com.

The project will develop predictive models for two questions:

- 1) Will the movie make money, lose money, or break even (approximately)?
- 2) What is the anticipated gross margin (profit) for the movie?

Data Exploration

Data Exploration

To this point we've removed the data columns for the variables that we will not be using in the analysis. The columns remaining in the data set are the following:

```
## [1] "duration"           "director_facebook_likes"
## [3] "actor_3_facebook_likes" "actor_1_facebook_likes"
## [5] "gross"              "movie_title"
## [7] "num_voted_users"    "cast_total_facebook_likes"
## [9] "facenumber_in_poster" "content_rating"
## [11] "budget"             "title_year"
## [13] "actor_2_facebook_likes" "imdb_score"
```

After exploring the data, we noticed there is a scattering of NAs across the variables. Due to the relatively low number of total NAs, we choose to remove all rows with NAs, leaving 3,828 rows of data.

Next we will explore the nature of the data for the variables we will be using in the analysis.

VAR	TYPE
duration	integer
director_facebook_likes	integer
actor_3_facebook_likes	integer
actor_1_facebook_likes	integer
gross	integer
movie_title	character
num_voted_users	integer
cast_total_facebook_likes	integer
facenumber_in_poster	integer

VAR	TYPE
content_rating	character
budget	double
title_year	integer
actor_2_facebook_likes	integer
imdb_score	double

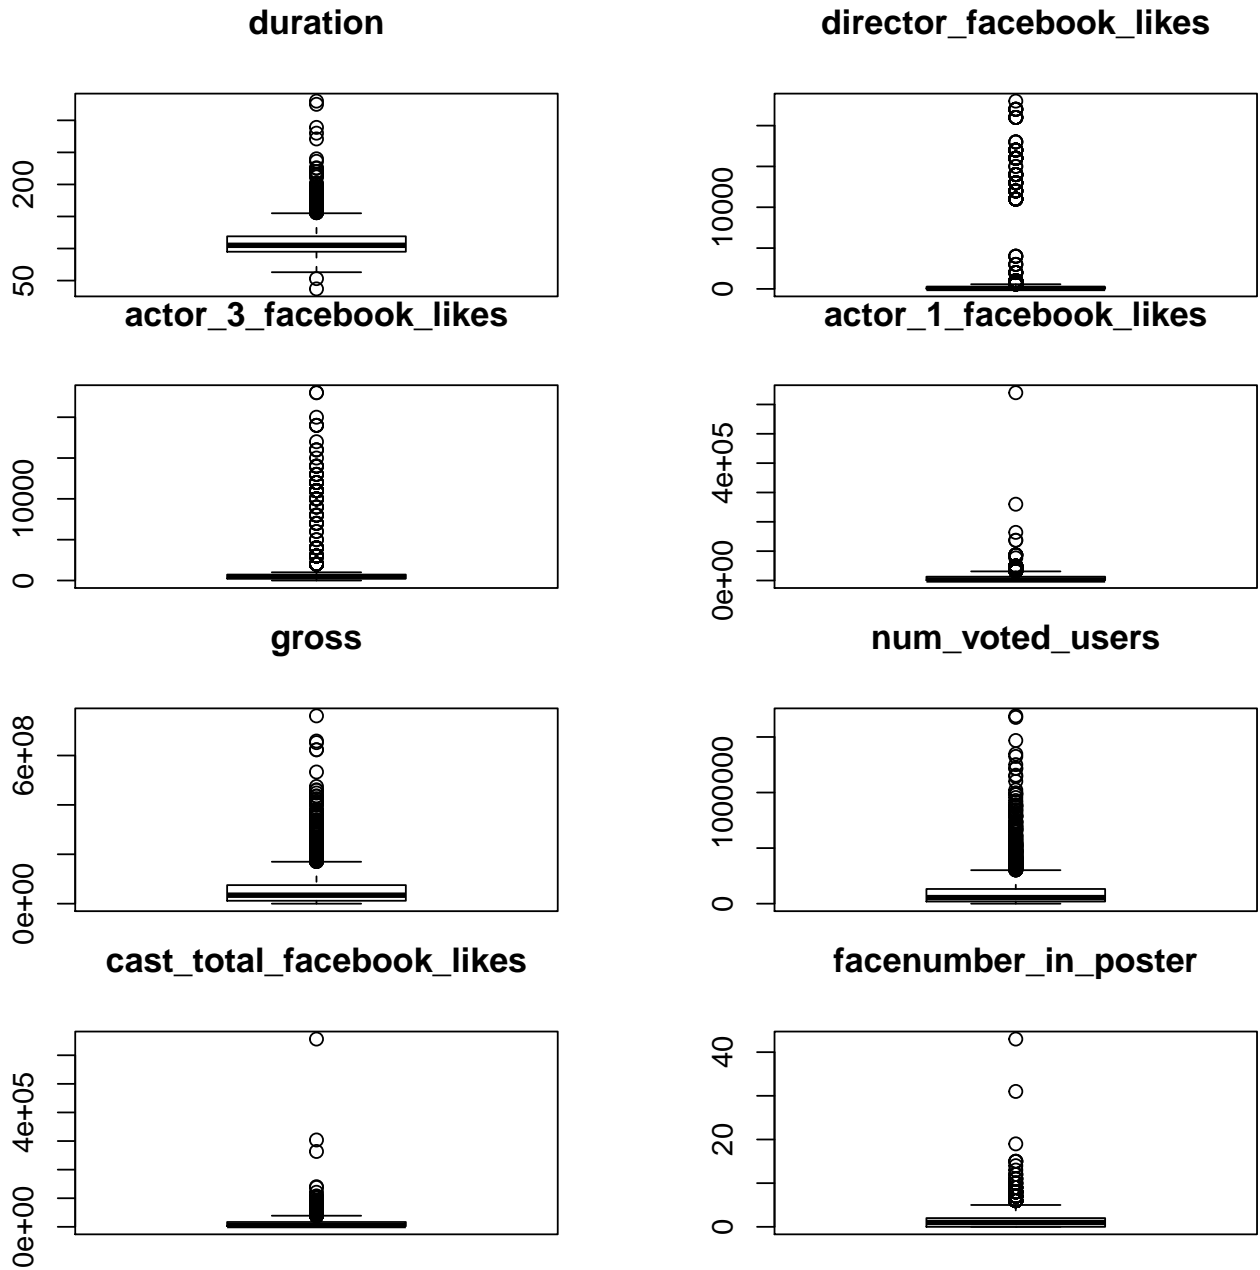
```
##      duration      director_facebook_likes actor_3_facebook_likes
## Min.   : 37.0    Min.   :    0.0          Min.   :    0.0
## 1st Qu.: 95.0    1st Qu.:   11.0          1st Qu.:  233.0
## Median :105.0    Median :   62.0          Median :  467.0
## Mean   :109.5    Mean   :  911.3          Mean   :  836.2
## 3rd Qu.:119.0    3rd Qu.:  235.0          3rd Qu.:  723.0
## Max.   :330.0    Max.   :23000.0          Max.   :23000.0
##
## actor_1_facebook_likes      gross      movie_title
## Min.   :    0.0      Min.   :    703  Length:3042
## 1st Qu.:   811.2      1st Qu.: 11787482  Class :character
## Median :  2000.0      Median : 34264376  Mode  :character
## Mean   :   8241.5      Mean   : 57651658
## 3rd Qu.: 13000.0      3rd Qu.: 75074326
## Max.   :640000.0      Max.   :760505847
##
## num_voted_users      cast_total_facebook_likes facenumber_in_poster
## Min.   :    22      Min.   :    0          Min.   : 0.000
## 1st Qu.: 19117      1st Qu.:  2210          1st Qu.: 0.000
## Median :  54462      Median :   4517          Median : 1.000
## Mean   : 108285      Mean   : 12340          Mean   : 1.419
## 3rd Qu.: 132124      3rd Qu.: 16904          3rd Qu.: 2.000
## Max.   :1689764      Max.   :656730          Max.   :43.000
##
##      content_rating      budget      title_year
## R      :1333      Min.   :    218      Min.   :1929
## PG-13   :1110      1st Qu.: 10725000      1st Qu.:1999
## PG      :  472      Median : 25000000      Median :2004
## G       :   70      Mean   : 40319361      Mean   :2003
## Not Rated:  18      3rd Qu.: 55000000      3rd Qu.:2010
## Unrated  :  13      Max.   :300000000      Max.   :2016
## (Other)  :  26
## actor_2_facebook_likes      imdb_score
## Min.   :    0.0      Min.   :1.600
## 1st Qu.:  436.0      1st Qu.:5.800
## Median :   729.5      Median :6.500
## Mean   :  2180.3      Mean   :6.383
## 3rd Qu.: 1000.0      3rd Qu.:7.100
## Max.   :137000.0      Max.   :9.300
##
```











	duration	director_facebook_likes	actor_3_facebook_likes	actor_1_facebook_likes
duration	1.0000000	0.2104197	0.1448777	0.0912903
director_facebook_likes	0.2104197	1.0000000	0.1219467	0.0868426
actor_3_facebook_likes	0.1448777	0.1219467	1.0000000	0.2483043
actor_1_facebook_likes	0.0912903	0.0868426	0.2483043	1.0000000
num_voted_users	0.3705768	0.3190331	0.2818195	0.1741973
cast_total_facebook_likes	0.1349956	0.1172865	0.4830033	0.9459350
facenumber_in_poster	0.0065845	-0.0523321	0.1042739	0.0538466
budget	0.2988689	0.0942904	0.2747815	0.1551897
title_year	-0.1086958	-0.0580504	0.1277213	0.0914452
actor_2_facebook_likes	0.1504159	0.1192872	0.5521997	0.3798140
imdb_score	0.3819342	0.2225461	0.0882029	0.1178984








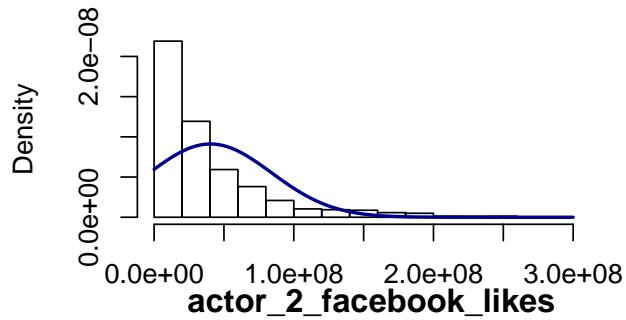
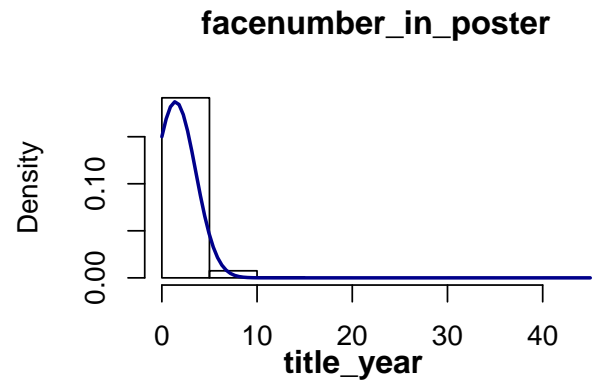

	num_voted_users	cast_total_facebook_likes	facenumber_in_poster	budget
duration	0.3705768	0.1349956	0.0065845	0.2988689
director_facebook_likes	0.3190331	0.1172865	-0.0523321	0.0942904
actor_3_facebook_likes	0.2818195	0.4830033	0.1042739	0.2747815
actor_1_facebook_likes	0.1741973	0.9459350	0.0538466	0.1551897
num_voted_users	1.0000000	0.2486828	-0.0441983	0.4054595
cast_total_facebook_likes	0.2486828	1.0000000	0.0750811	0.2362870
facenumber_in_poster	-0.0441983	0.0750811	1.0000000	-0.0267742
budget	0.4054595	0.2362870	-0.0267742	1.0000000
title_year	0.0241674	0.1256809	0.0873375	0.2412454
actor_2_facebook_likes	0.2524944	0.6319688	0.0625703	0.2526741
imdb_score	0.5089320	0.1377072	-0.0694804	0.0713682

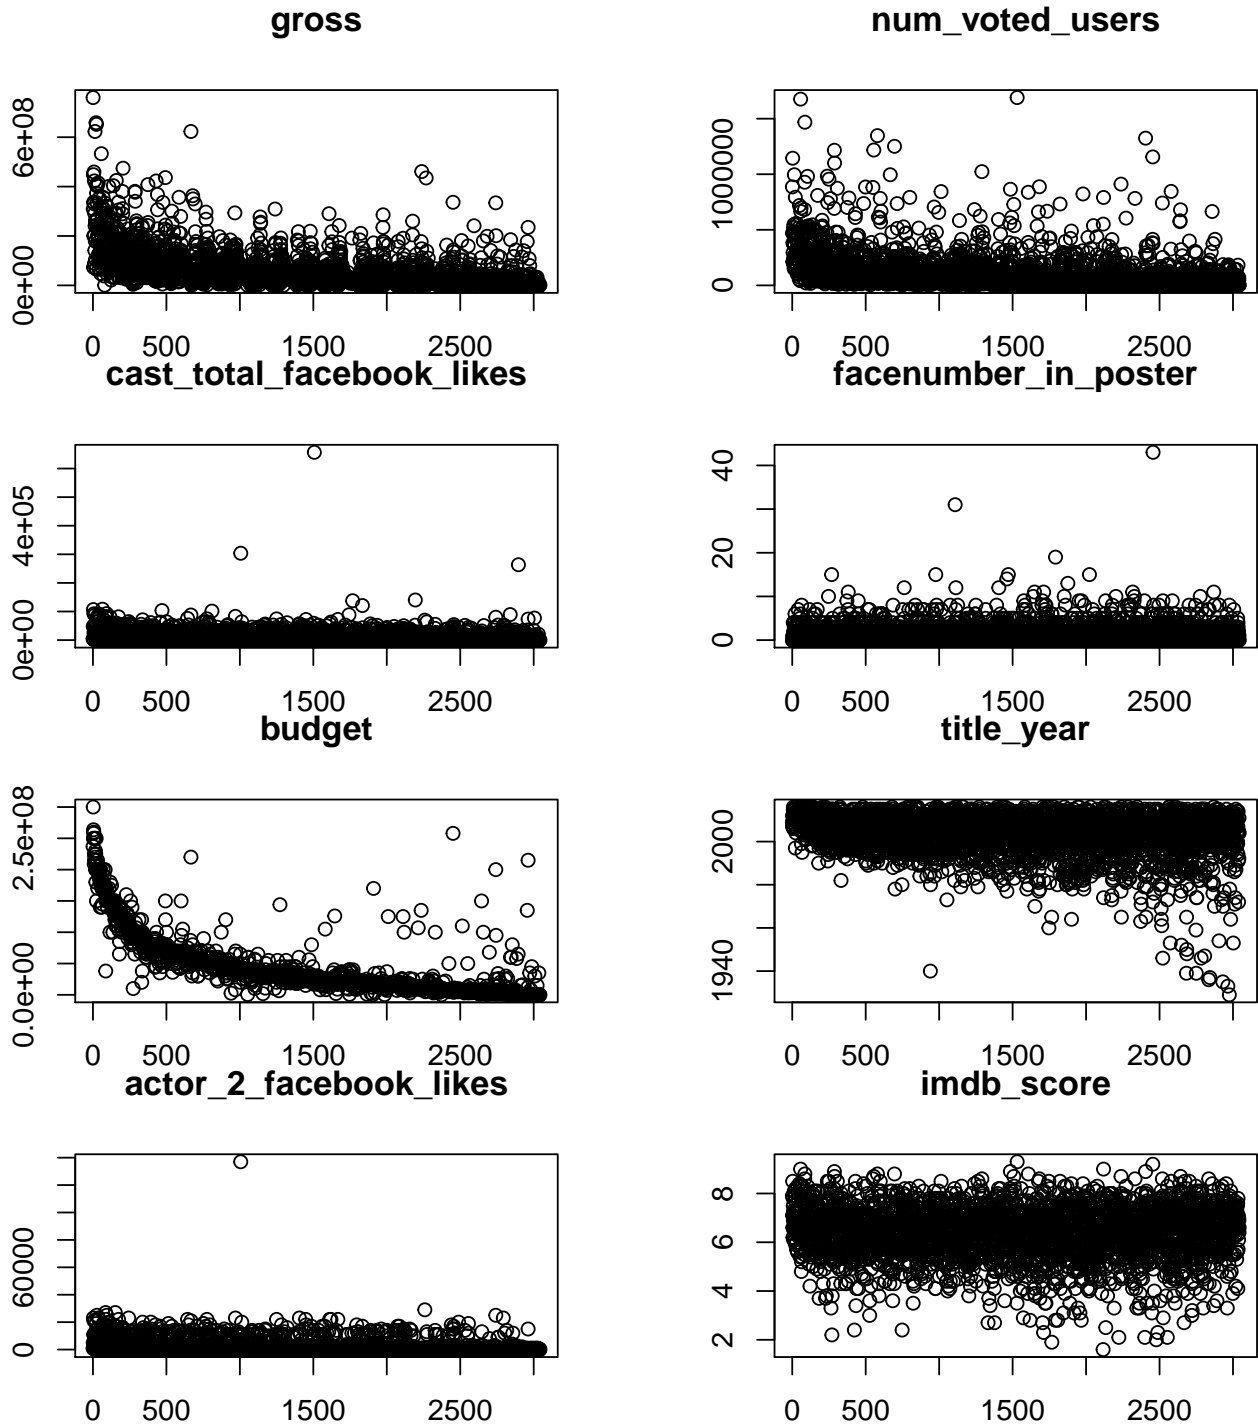
	title_year	actor_2_facebook_likes
duration	-0.1086958	0.1504159
director_facebook_likes	-0.0580504	0.1192872
actor_3_facebook_likes	0.1277213	0.5521997
actor_1_facebook_likes	0.0914452	0.3798140
num_voted_users	0.0241674	0.2524944
cast_total_facebook_likes	0.1256809	0.6319688
facenumber_in_poster	0.0873375	0.0625703
budget	0.2412454	0.2526741
title_year	1.0000000	0.1253783

	title_year	actor_2_facebook_likes
actor_2_facebook_likes	0.1253783	1.0000000
imdb_score	-0.1504498	0.1274387









As we can see from the plots and statistical summary, most of the variables have a reasonable distribution except those variable associated with the Facebook likes. There are five variables related to Facebook likes that are highly skewed due to a large number of zeros. At this point we assume these zeros represent NAs in the Facebook data.

Next, we'll use the mice package to impute the Facebook likes data for the zeros/NAs.

```
##      actor_1_facebook_likes cast_total_facebook_likes
## 2502                      1                      1
```

```

## 520          1          1
## 10          1          1
## 1          1          1
## 6          1          1
## 2          1          1
## 1          0          0
##          1          1
##      actor_2_facebook_likes actor_3_facebook_likes director_facebook_likes
## 2502          1          1          1
## 520          1          1          0
## 10          1          0          1
## 1          1          0          0
## 6          0          0          1
## 2          0          0          0
## 1          0          0          1
##          9          20          523
##
## 2502  0
## 520  1
## 10  1
## 1  2
## 6  2
## 2  3
## 1  4
##      554

```

```

##      duration      director_facebook_likes actor_3_facebook_likes
## Min.   : 37.0    Min.   : 2          Min.   : 2.0
## 1st Qu.: 95.0    1st Qu.: 30          1st Qu.: 233.0
## Median :105.0    Median : 98          Median : 467.5
## Mean   :109.5    Mean   : 1107         Mean   : 840.9
## 3rd Qu.:119.0    3rd Qu.: 295         3rd Qu.: 723.0
## Max.   :330.0    Max.   :23000        Max.   :23000.0
##

```

```

## actor_1_facebook_likes      gross      movie_title
## Min.   : 2          Min.   : 703      Length:3042
## 1st Qu.: 812          1st Qu.: 11787482      Class :character
## Median : 2000          Median : 34264376      Mode  :character
## Mean   : 8242          Mean   : 57651658
## 3rd Qu.: 13000          3rd Qu.: 75074326
## Max.   :640000          Max.   :760505847
##

```

```

## num_voted_users      cast_total_facebook_likes facenumber_in_poster
## Min.   : 22          Min.   : 2          Min.   : 0.000
## 1st Qu.: 19117          1st Qu.: 2212          1st Qu.: 0.000
## Median : 54462          Median : 4523          Median : 1.000
## Mean   : 108285          Mean   : 12348          Mean   : 1.419
## 3rd Qu.: 132124          3rd Qu.: 16919          3rd Qu.: 2.000
## Max.   :1689764          Max.   :656730          Max.   :43.000
##

```

```

##      content_rating      budget      title_year
## R      :1333    Min.   : 218    Min.   :1929
## PG-13   :1110    1st Qu.: 10725000    1st Qu.:1999
## PG      : 472    Median : 25000000    Median :2004

```




```
## G      : 70   Mean   : 40319361   Mean    :2003
## Not Rated: 18   3rd Qu.: 55000000   3rd Qu.:2010
## Unrated  : 13   Max.    :300000000   Max.    :2016
## (Other)  : 26
## actor_2_facebook_likes   imdb_score
## Min.    :      2.0       Min.    :1.600
## 1st Qu.:   436.2       1st Qu.:5.800
## Median  :   730.0       Median  :6.500
## Mean    :  2180.6       Mean    :6.383
## 3rd Qu.:  1000.0       3rd Qu.:7.100
## Max.    :137000.0       Max.    :9.300
##
```

Data Preparation

Data Preparation


One of the big issues faced in when using this dataset is the time frame. These movies were collected over the past 80+ years, and the following shows our distribution over time:




As you can see, the vast majority came from 1990s and above, but we can't discredit the movies from previous year. In order to accurately portray elements from the past, we have instituted a rate of inflation calculation. Using the consumer price index (for our part here we are making a crucial assumption, that all dollars are calculated based on US currency, and we are ignoring even more complex foreign exchange rates of the time), we can calculate the gross value per year. As a basis of comparison, we are using the CPI index from 2016, as the last movie was made in 2016.

```
## The following object is masked _by_ .GlobalEnv:
##
## cpi
```

Unadjusted Gross Per Year




Adjusted Gross Per Year



From the above graphs, we can see that the adjustment for the gross did indeed create a more uniform dataset (where as before we saw movies increasing over the years). As a point of interest, the movies that made over a billion dollars are shown below:

##	movie_title	gross	adj_gross
## 5	Snow White and the Seven Dwarfs	184925485	3082091417
## 7	Gone with the Wind	198655278	3430019188
## 8	Pinocchio	84300000	1445142857
## 26	The Sound of Music	163214286	1243537417
## 39	The Exorcist	204565000	1105756757
## 48	Jaws	260000000	1159851301
## 53	Star Wars: Episode IV - A New Hope	460935665	1825487782
## 90	E.T. the Extra-Terrestrial	434949459	1081739587

A quick Google search indicates that the above movies are consistently listed the top grossing movies of all time. Furthermore, our “estimated adjusted gross” mimics the findings that we see with adjusted gross (for the most part, there are two schools of thought on how to adjust gross, using ticket prices or our method adjusting based on CPI). Though our dollar amount vary slightly from other sources, any variance is consistent across our dataset.



Build Models

Build Models

Binomial Regression

Our first model we want to investigate is whether or not we can predict if film will make money given the cast and direction. To do this, we decided to create a binary regression model, transforming our adjusted

```
##
## Call:
## glm(formula = money ~ ., family = binomial(link = "logit"), data = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -4.5060  -1.1107   0.5122   1.0648   1.8662
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)      8.006e+01  1.066e+01   7.513 5.79e-14 ***
## title_year       -3.950e-02  5.302e-03  -7.450 9.34e-14 ***
## duration         -1.327e-02  2.471e-03  -5.370 7.86e-08 ***
## director_facebook_likes -3.318e-05  1.433e-05  -2.315  0.0206 *
## actor_3_facebook_likes -1.248e-04  7.485e-05  -1.668  0.0954 .
## actor_1_facebook_likes -1.222e-04  5.038e-05  -2.425  0.0153 *
## num_voted_users      8.635e-06  6.478e-07  13.328 < 2e-16 ***
## cast_total_facebook_likes 1.179e-04  5.034e-05   2.343  0.0191 *
## facenumber_in_poster  4.012e-02  2.223e-02   1.805  0.0710 .
## actor_2_facebook_likes -1.233e-04  5.259e-05  -2.345  0.0190 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
##      Null deviance: 3307.1  on 2432  degrees of freedom
```

```
## Residual deviance: 2957.1 on 2423 degrees of freedom
## AIC: 2977.1
##
## Number of Fisher Scoring iterations: 5


## [1] 0.7446781
```

Using all the prediction variables at hand, the model accurately predicts 76% of the time. Using backward stepwise regression, we attempted to remove some variables that may not have had significance in our model.

```
## Start: AIC=2977.14
## money ~ title_year + duration + director_facebook_likes + actor_3_facebook_likes +
## actor_1_facebook_likes + num_voted_users + cast_total_facebook_likes +
## facenumber_in_poster + actor_2_facebook_likes
##
##              Df Deviance    AIC
## <none>              2957.1 2977.1
## - actor_3_facebook_likes    1   2959.9 2977.9
## - facenumber_in_poster      1   2960.5 2978.5
## - director_facebook_likes   1   2962.5 2980.5
## - actor_2_facebook_likes    1   2962.9 2980.9
## - cast_total_facebook_likes 1   2962.9 2980.9
## - actor_1_facebook_likes    1   2963.3 2981.3
## - duration                  1   2987.1 3005.1
## - title_year                1   3019.6 3037.6
## - num_voted_users           1   3242.0 3260.0

##
## Call:
## glm(formula = money ~ title_year + duration + director_facebook_likes +
## actor_3_facebook_likes + actor_1_facebook_likes + num_voted_users +
## cast_total_facebook_likes + facenumber_in_poster + actor_2_facebook_likes,
## family = binomial(link = "logit"), data = train)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -4.5060  -1.1107   0.5122   1.0648   1.8662
##
## Coefficients:
##              Estimate Std. Error z value Pr(>|z|)
## (Intercept)    8.006e+01  1.066e+01   7.513 5.79e-14 ***
## title_year     -3.950e-02  5.302e-03  -7.450 9.34e-14 ***
## duration       -1.327e-02  2.471e-03  -5.370 7.86e-08 ***
## director_facebook_likes -3.318e-05  1.433e-05  -2.315  0.0206 *
## actor_3_facebook_likes -1.248e-04  7.485e-05  -1.668  0.0954 .
## actor_1_facebook_likes -1.222e-04  5.038e-05  -2.425  0.0153 *
## num_voted_users  8.635e-06  6.478e-07  13.328 < 2e-16 ***
## cast_total_facebook_likes 1.179e-04  5.034e-05   2.343  0.0191 *
## facenumber_in_poster  4.012e-02  2.223e-02   1.805  0.0710 .
## actor_2_facebook_likes -1.233e-04  5.259e-05  -2.345  0.0190 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```

```
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 3307.1 on 2432 degrees of freedom
## Residual deviance: 2957.1 on 2423 degrees of freedom
## AIC: 2977.1
##
## Number of Fisher Scoring iterations: 5
```



```
## [1] 0.7446781
```

Profit Margin Model

```
##          title_year          duration
##              0              0
## director_facebook_likes actor_3_facebook_likes
##              0              0
## actor_1_facebook_likes          gross
##              0              0
## num_voted_users cast_total_facebook_likes
##              0              0
## facenumber_in_poster          budget
##              0              0
## actor_2_facebook_likes          imdb_score
##              0              0
##              cpi          adj_gross
##              0              0
## adj_budget          adj_margin
##              0              0

##          title_year          duration
```

```





##              0              0
## director_facebook_likes actor_3_facebook_likes
##              0              0
## actor_1_facebook_likes          gross
##              0              0
## num_voted_users cast_total_facebook_likes
##              0              0
## facenumber_in_poster          budget
##              1259              0
## actor_2_facebook_likes          imdb_score
##              0              0
## cpi          adj_gross
##              0              0
## adj_budget          adj_margin
##              0              0

##
## Call:
## lm(formula = adj_gross ~ (duration + director_facebook_likes +
## actor_3_facebook_likes + actor_1_facebook_likes + num_voted_users +
## cast_total_facebook_likes + actor_2_facebook_likes + imdb_score +
## adj_budget + adj_margin) - adj_margin, data = movies1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -522395904 -36484452 -14594206  15010635 3250078902
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    -8.810e+07  1.713e+07  -5.143 2.88e-07 ***
## duration         2.469e+05  1.167e+05   2.116 0.034451 *
## director_facebook_likes -1.482e+03  6.491e+02  -2.284 0.022459 *
## actor_3_facebook_likes -1.023e+04  3.153e+03  -3.245 0.001188 **
## actor_1_facebook_likes -7.661e+03  1.924e+03  -3.981 7.03e-05 ***
## num_voted_users    2.811e+02  1.878e+01  14.962 < 2e-16 ***
## cast_total_facebook_likes 7.338e+03  1.921e+03   3.820 0.000136 ***
## actor_2_facebook_likes -8.032e+03  2.033e+03  -3.951 7.96e-05 ***
## imdb_score        1.239e+07  2.588e+06   4.787 1.78e-06 ***
## adj_budget         6.723e-01  5.128e-02  13.110 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 121600000 on 3032 degrees of freedom
## Multiple R-squared:  0.2681, Adjusted R-squared:  0.266
## F-statistic: 123.4 on 9 and 3032 DF, p-value: < 2.2e-16

## [1] 121906.9



## 'log Lik.' -60942.45 (df=11)

```



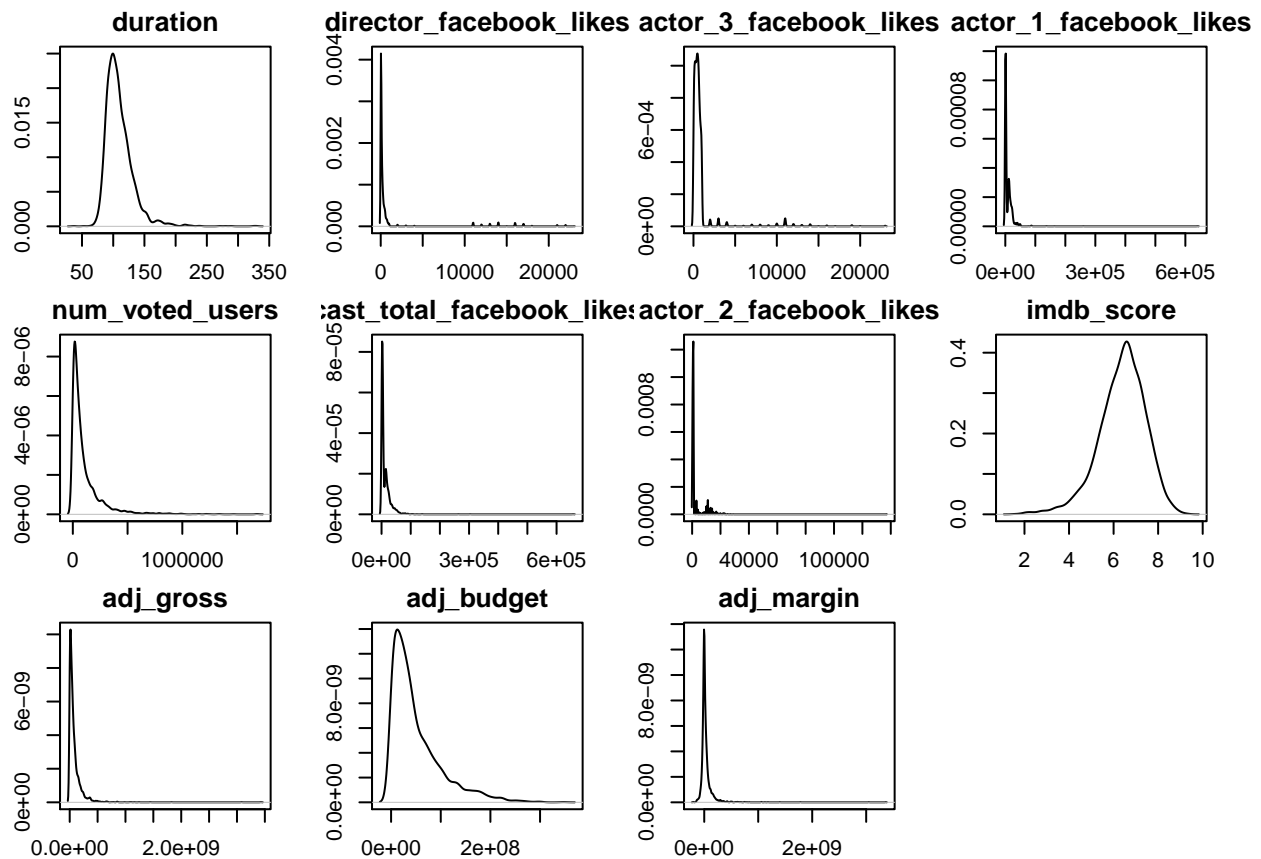

Histogram of movies1\$adj_gross

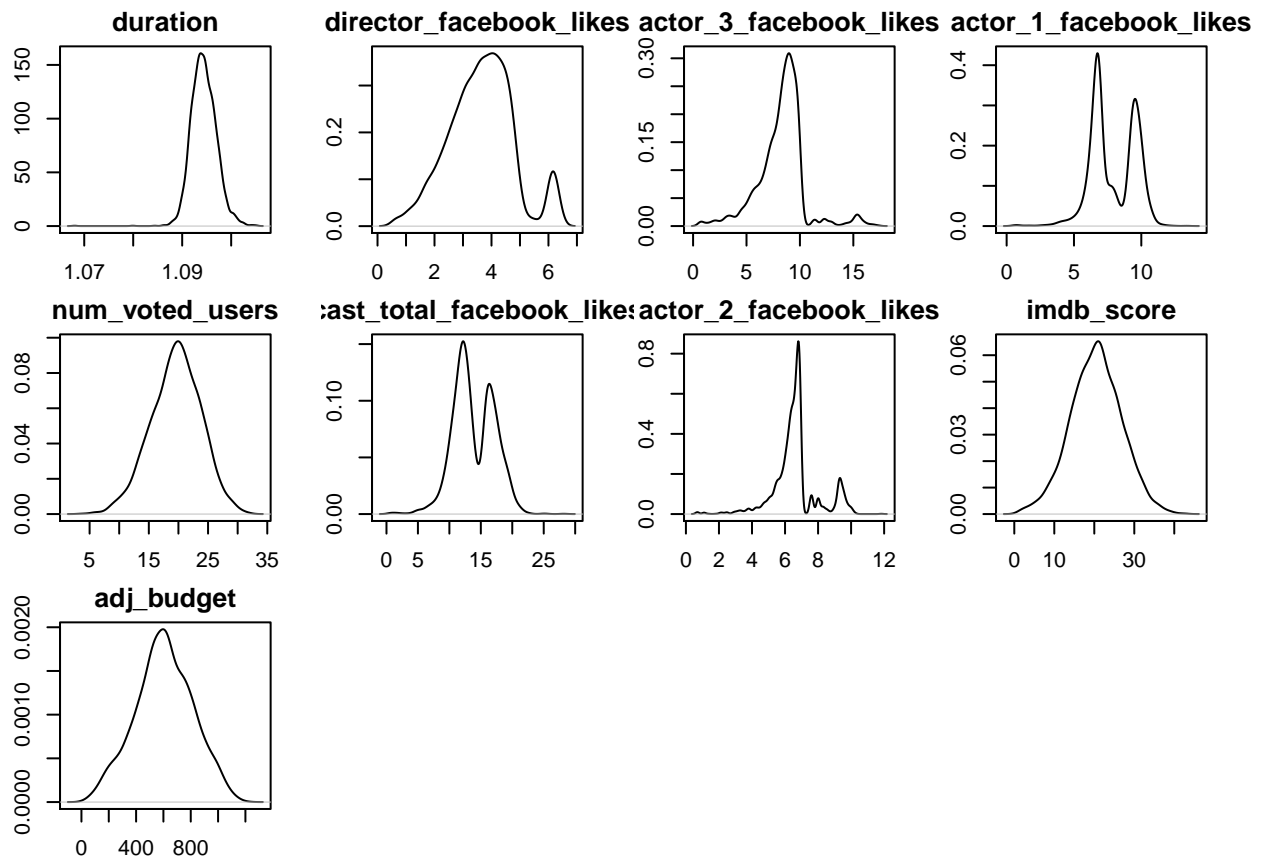
Histogram of fitted(m1_back)







```
## [1] 13.64336
## attr(,"method")
## [1] "moment"
```

Histogram of m1_back\$residuals






```
##
## Call:
## lm(formula = adj_gross ~ actor_1_facebook_likes + num_voted_users +
##     cast_total_facebook_likes + actor_2_facebook_likes + imdb_score +
##     adj_budget, data = movies2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -213152575  -48916804  -14870946   25855572  3210858733
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   -158397352    16543217  -9.575  < 2e-16 ***
## actor_1_facebook_likes   -36912913     6450144  -5.723  1.15e-08 ***
## num_voted_users         8837756      796101  11.101  < 2e-16 ***
## cast_total_facebook_likes  15144212    3928800   3.855  0.000118 ***
## actor_2_facebook_likes   -4672309     3280199  -1.424  0.154434
## imdb_score          3753304      429442   8.740  < 2e-16 ***
## adj_budget           170799       13378  12.767  < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.23e+08 on 3035 degrees of freedom
## Multiple R-squared:  0.2506, Adjusted R-squared:  0.2491
## F-statistic: 169.2 on 6 and 3035 DF, p-value: < 2.2e-16
```




```
## [1] 121973
```

```
## 'log Lik.' -60978.48 (df=8)
```

Histogram of movies2\$adj_gross




Histogram of fitted(m2_back)



```
## [1] 12.63913
## attr("method")
## [1] "moment"
```

Histogram of m2_back\$residuals



##	Movie Title	Actual	Adjusted Gross	Predicted Gross
## 1	The Broadway Melody		39181395	4377823
## 2	42nd Street		42790698	38466370
## 3	Top Hat		52554745	50538646
## 4	Modern Times		2818619	174017497
## 5	Snow White and the Seven Dwarfs		3082091417	181193243
## 6	The Wizard of Oz		383354452	210022805

##	Actual Profit Margin	Predicted Profit Margin
## 1	0.8650285	-0.2079914
## 2	0.8091304	0.7876732
## 3	0.7970000	0.7889019
## 4	-8.1886428	0.8511685
## 5	0.9891848	0.8160343
## 6	0.8738887	0.7698092

This line of code just to prove that, there is no correlation between the quality of the moview with

```
##
## Pearson's product-moment correlation
##
## data: movies$imdb_score and movies$profit_margin
## t = 2.6579, df = 3040, p-value = 0.007905
## alternative hypothesis: true correlation is not equal to 0
## 95 percent confidence interval:
## 0.01263244 0.08354498
## sample estimates:
## cor
## 0.04814938
```

Smooth Operators - All Done!